

AForce OS — Engineering Brief

AForce OS — Engineering Brief for Provisional Patent Application

Prepared for: Outside counsel, AForce IP filing **Application:** AForce OS mobile application (Expo / React Native) **Date:** April 29, 2026 **Companion artifacts:** Patent figures FIG. 1 – FIG. 10
(attached_assets/patent/fig{1..10}-*.svg)

This brief documents what AForce engineering has built and shipped in the AForce OS mobile application as of the date above. Every system below is implemented in production code, not vaporware. Each section cites exact file paths and line numbers in the repository so counsel can drop them directly into the specification or hand them to a technical reviewer.

The five systems counsel asked about — **Mobile Application**, **Scoring Engine**, **Deterministic AI Layer**, **Autopilot**, and **Autoscan** — are described in that order. Each maps to one of the attached patent figures (noted at the top of each section). Wireframes referenced throughout are also live, interactive iframes on the engineering canvas (AForce OS app preview, Recovery Card unlocked/locked, Hydration Score Band, Hero band).

1. Mobile Application

Maps to: FIG. 1 (System Architecture), FIG. 7 (Mobile App Architecture), FIG. 6 (Wireframes)
Repository root: artifacts/aforce-os/ **Stack:** Expo (React Native), expo-router for file-based navigation, Clerk for auth, Postgres + Drizzle on the backend (artifacts/api-server/), WebSocket for multi-device sync.




FIG. 1 — System Architecture.

1.1 Navigation skeleton

File-based routing through `expo-router`. Root `Stack` wraps a persistent bottom `Tabs` layout for the main experience.

Route	File	Purpose
/ (tab)	app/(tabs)/index.tsx	Hydration Control Center — animated status orb, AI command, prediction strip
/check (tab)	app/(tabs)/check.tsx	Performance Signals — biometric snapshot
/protocol (tab)	app/(tabs)/protocol.tsx	Active recovery protocols
/store (tab)	app/(tabs)/store.tsx	Product shopping (RTD, sticks, canister)
/profile (tab)	app/(tabs)/profile.tsx	User settings, HealthKit integration, Protocol Tools
/scan (modal)	app/scan.tsx	Camera-based product scanning (Autoscan)
/heat	app/heat.tsx	Heat-risk prediction dashboard
/sweat	app/sweat.tsx → screens/SweatCalculatorScreen.tsx	Sweat Intelligence v2 calculator (drives Autopilot)
/competition	app/competition.tsx	GPS-based territory map
/subscription	app/subscription.tsx	Premium plan management
/(auth)/sign-in , /(auth)/sign-up	app/(auth)/*	Clerk-backed auth

1.2 State management




FIG. 7 — Mobile App Architecture.

Pattern: Sliced React Context + `useReducer`. The monolithic `AppState` is projected through memoized slice contexts so unrelated UI sections do not re-render on unrelated state changes.

Layer	File	Responsibility
-------	------	----------------

Layer	File	Responsibility
Store orchestrator	store/useAppStore.tsx	Owns AppState , dispatch, periodic 30s server polling, WebSocket subscription
Reducer (pure)	store/appStoreReducer.ts	All state transitions: intake logging, timer ticks, autopilot, engine refreshes
Slice projections	store/slices.tsx	EngineSlice , UserSlice , SocialSlice , CycleSlice , InventorySlice , SweatAutopilotSlice
Type contract	store/appStoreTypes.ts , types/index.ts	AppState , UserState , ScoreEngineOutput , IntakeEvent , InventoryState , SweatAutopilot

1.3 Server / data layer

Component	File	Notes
API client	lib/api.ts	Typed fetch wrapper, persists x-device-id for anonymous-then-merged sessions
Real API service	services/realApi.ts	Normalizes Postgres rows into UserState and re-runs the client-side scoring engine on every response so UI always shows a freshly computed score
Realtime channel	services/realApi.ts#subscribeToStateUpdates	WebSocket push from api-server for multi-device sync
Local persistence	AsyncStorage	deviceId , Clerk tokenCache , aforce.subscription flag

1.4 Hardware integrations

Surface	File	Notes
Voice	services/voiceService.ts + voiceTemplateEngine.ts	Classifies transcripts into intents (LOG_INTAKE , GET_STATUS , etc.) and renders brand-aligned spoken responses
Camera	services/hydrationScanService.ts , components/CameraScanModal.tsx	expo-camera CameraView reading ean13 , upc_a , qr , code128
Apple HealthKit	services/appleHealth.ts	restingHeartRate , hrvSdnn , sleepHoursLastNight fed into the scoring engine
Auth	components/ClerkAuthBridge.tsx	Bridges Clerk session → internal store; (tabs) layout redirects unauthenticated users
Entitlement	hooks/useEntitlement.ts	Feature gating (subscriber vs. free)

1.5 Wireframes (live on canvas)




FIG. 6 — Wireframes.

Counsel can view the actual rendered screens — not static mockups — in the engineering canvas:


Canvas shape ID	What it shows	Source
artifact: .../aforce-os	Full live AForce OS app	artifacts/aforce-os/
recovery-unlocked	Recovery Card — subscriber view	artifacts/mockup-sandbox/src/components/mockups/recovery-mode/Unlocked.tsx
recovery-locked	Recovery Card — paypal	.../recovery-mode/Locked.tsx
score-band	Hydration Score Band component	.../score-band/Default.tsx
score-band-hero	Hydration Score Band — hero variant	.../score-band/Hero.tsx

2. Scoring Engine

Maps to: FIG. 8 (Scoring Engine) **Primary file:** artifacts/aforce-os/utls/scoringEngine.ts **Supporting files:** services/hydrationScoreService.ts , types/index.ts , store/slices.tsx#useEngineSlice

FIG. 8 — Scoring Engine: Deterministic Pipeline

Pure function: same inputs → same outputs, no model state, no randomness



AForce OS · Patent Figure

FIG. 8 — Scoring Engine.

The Scoring Engine is a **pure function** that takes a `UserState` snapshot and returns a `ScoreEngineOutput`. It runs both server-side (during state writes) and client-side (on every API response, for sub-network-latency UI updates). Because it is pure, the same input snapshot always produces the same output — this is the foundation of the determinism claim in §3.

2.1 Inputs (`UserState` , `types/index.ts`)

- `intakeEvents[]` — recent logs with pre-computed absorption impact
- `ozConsumedToday` , `ozTarget` — base hydration progress
- `aforceUnitsToday` — count of AForce-specific intakes (drives the bonus term)
- `lastIntakeTime` — timestamp for decay and recovery-momentum calculations
- `heatLoad` , `sweatRate` , `activityLevel` — environmental + physiological stressors (0-10)
- `urineSignal` — biomarker proxy (1-8 scale)
- `symptomState` , `symptoms[]` — severity bucket + specific symptom IDs
- `appleHealth` — `hrvSdnn` , `sleepHoursLastNight` from HealthKit
- `clutchActive` — boolean; multiplies decay during high-intensity windows
- `socialMode` — alcohol mitigation state (drinks, ABV, food)

2.2 Output (`ScoreEngineOutput` , `types/index.ts:256`)

```

{
  score: number, // 0-100
  performanceState: { level, color, pulseSpeed },
  pulseConfig: { waveBehavior, animations },
  breakdown: ScoreContribution[], // {id, label, delta, maxMagnitude, hint}
  prediction: ScorePrediction, // {decayPerMinute, minutesToDepleted, label}
  command: AICommand, // {action, explanation, urgencyLevel, estimatedImpact}
  social: SocialRollup // BAC, impairment level, safety prompts
}

```

2.3 Score formula (`scoringEngine.ts:363`)

```

score = base_intake
+ aforce_bonus
+ recency
+ consistency
+ context
+ recovery_momentum
+ symptom_penalty      (negative)
+ urine_penalty        (negative)
+ output_stress        (negative)
+ sleep_carry
+ recovery_signal
+ confirmation

```

Each term is a bounded contribution with its own clamp and weight, exposed in the `breakdown` array so the UI can render an audit (the "Why this score?" sheet in `components/ScoreBreakdownSheet.tsx`).

2.4 Absorption model (`hydrationScoreService.ts:25`)

- **Plain water:** 60 % immediate, 40 % over 12.5 min
- **AForce:** 70 % immediate, 30 % over 25 min (faster onset because of structured-water + mineral pairing)
- **Cap:** at most 1.5 units per rolling 20-min window; excess is 75 % efficient (models gut absorption ceiling)

2.5 Continuous decay (`scoringEngine.ts:166-202`)

```

BaseDecay = 0.4 × (weight_lb / 150) + 0.1 × activity_level
+ heatModifier(temp_C > 25)
+ humidityModifier(RH > 50%)
× (clutchActive ? 1.3 : 1.0)

```

Decay is applied per-minute and surfaces in the UI as the prediction string ("Drops to DEPLETED in 1 h 12 m").

2.6 Score bands (`scoringEngine.ts:38-42`)

Band	Range	Default recheck interval (<code>scoringEngine.ts:444-447</code>)
PEAK	90 - 100	≥ 20 min
BALANCED	75 - 89	≥ 15 min
RECOVERING	60 - 74	≥ 10 min
DEPLETED	0 - 59	≥ 5 min

These bands are overridden by the Autopilot system (§4) when an active sweat session is in progress.

2.7 UI consumption

Driven exclusively by `useEngineSlice()` (`store/slices.tsx`). Primary consumers:

- `components/StatusPulseOrb.tsx` — uses `pulseConfig` for the animated orb
- `components/AnimatedScore.tsx` — counts the score up/down
- `components/ScoreBreakdownSheet.tsx` — renders `breakdown[]`
- `components/home/OrbSection.tsx` — renders `command.action` + `command.explanation`
- Home screen "Prediction Strip" — renders `prediction.label`

3. Deterministic AI Layer

Maps to: FIG. 9 (Deterministic AI Layer) **Primary file:** `artifacts/aforce-os/utls/scoringEngine.ts` (command generator at lines 537-591) **Supporting:** `services/bacEstimationService.ts`, `services/legalSafetyService.ts`




FIG. 9 — Deterministic AI Layer.

3.1 What "Deterministic" means here

Conventional consumer health apps make recommendations through probabilistic LLM calls. The AForce Deterministic AI Layer is the opposite: an **ordered, pure-function rule pipeline** with **no remote inference and no nondeterministic LLM calls** in the safety- and performance-critical path. The same `UserState` snapshot always produces the same `AICommand`. This matters for three reasons that the application leans on as differentiators:

1. **Auditability.** Every output can be reproduced offline by replaying the snapshot through the same code. The `breakdown[]` array exposes every contributing term.
2. **Safety-class behavior.** Alcohol-mitigation recommendations ("DO NOT DRIVE", "PULL FROM ROTATION") cannot be subject to model drift or hallucination.
3. **Edge-first.** Decisions are computed locally on the device (or in the same reducer function on the server). The app degrades gracefully offline.

3.2 Rule pipeline (Stage 3)

The pipeline (FIG. 9) runs in this order, top-down, on each engine refresh:

1. **Mode classification** — Is the user in a *Social Mode* override (alcohol present)? If yes, skip the standard ladder; the Social Mode rule set takes precedence (`scoringEngine.ts:458-531`). Otherwise, use the standard PEAK / BALANCED / RECOVERING / DEPLETED ladder from §2.6.
2. **Penalty application** — Symptom severity, urine-color signal, output stress, and HRV/sleep deficits each apply a bounded negative delta.
3. **Bonus application** — AForce intakes, consistency streaks, recovery momentum, and confirmed intakes each apply a bounded positive delta.
4. **Decay projection** — Continuous decay model (§2.5) yields `prediction.minutesToDepleted` and the prediction label.
5. **Command generation** — `generateCommand()` (lines 537-591) selects a "WHAT + WHEN + OUTCOME" directive based on state, deficit, and inventory.
6. **Recheck cadence** — Default cadence by band (§2.6); if Autopilot is active, replaced by the autopilot interval (§4.3).

3.3 Social Mode override (scoringEngine.ts:458-531)




FIG. 4 — Social Mode Flow.

When `socialMode.active === true` , the engine swaps in a parallel ruleset that:

- Estimates BAC via the **Widmark approximation** (services/bacEstimationService.ts:92-147), parameterized by sex, weight, drinks consumed, ABV, time elapsed, and food intake.
- Maps BAC to one of five **impairment levels** — LOW , ELEVATED , MODERATE , HIGH , CRITICAL (services/legalSafetyService.ts:28-37).
- Issues escalating commands and surfaces transportation / legal safety language. The disclaimer copy is locked and globally safe.

3.4 Guardian Risk Engine and Clutch Strategy

Two specialized branches sit on top of the core pipeline:

- **Guardian Risk Engine** (scoringEngine.ts:728-748) — applies an additional risk overlay when symptoms + heat + low HRV co-occur, escalating the command urgency and tightening the recheck cadence.
- **Clutch Strategy** (scoringEngine.ts:790-838) — when `clutchActive` is set (user has marked a high-intensity competitive window), the engine multiplies decay by 1.3, prefers AForce-led commands, and shortens cadence.

3.5 Outputs

- `command.action` — short imperative ("16 oz + 2 sticks now")
- `command.explanation` — rationale shown beneath the command
- `command.urgencyLevel` — low | medium | high | critical
- `command.estimatedImpact` — predicted score delta if executed
- `pulseConfig` — drives the orb animation so the user can read state at a glance
- `recheckMinutes` — fed to the timer (overridden by Autopilot when applicable)

4. Autopilot

Maps to: FIG. 10 (Autopilot / Autoscan) **Engine file:** artifacts/aforce-os/services/sweatRateEngine.ts **Wiring:** screens/SweatCalculatorScreen.tsx , store/appStoreReducer.ts , hooks/useHeatGuard.ts , components/RiskTimerDisplay.tsx




FIG. 10 — Autopilot / Autoscan.

4.1 What Autopilot is

Autopilot is a **physiologically-paced recheck cadence** that takes over the timer for a fixed window (4 h) following a measured sweat session. Instead of using the band-default cadence from §2.6, the timer is driven by the user's just-measured **hydration deficit %** so high-deficit athletes get checked on more aggressively for the duration of the recovery window.

4.2 Trigger

The user runs the Sweat Calculator in any of three modes (Quick / Precision / Estimate) and taps Calculate. `commitSession()` (SweatCalculatorScreen.tsx:167-172) dispatches `SET_SWEAT_AUTOPILOT` with the session's autopilot derivation.

4.3 Cadence ladder (sweatRateEngine.ts:489-493 , verified)

```
export function deriveAutopilot(deficitPct: number): SweatAutopilot {  
  if (deficitPct >= 4) return { intervalMin: 8, urgency: 'critical', recoveryWindowHours: 4 };  
  if (deficitPct >= 2) return { intervalMin: 12, urgency: 'high', recoveryWindowHours: 4 };  
  return { intervalMin: 20, urgency: 'moderate', recoveryWindowHours: 4 };  
}
```

Deficit	Interval	Urgency	Window
≥ 4 %	8 min	critical	4 h
≥ 2 %	12 min	high	4 h

Deficit	Interval	Urgency	Window
< 2 %	20 min	moderate	4 h

4.4 Reducer wiring (store/appStoreReducer.ts)

SET_SWEAT_AUTOPILOT does three things atomically:

1. Stores `sweatAutopilot` (interval + urgency) and `sweatAutopilotSetAt` (timestamp, used to enforce the 4 h window).
2. **Resets** `state.timerSeconds = autopilot.intervalMin × 60` so the visible countdown immediately reflects the new cadence.
3. Clears any `pendingConfirmation` so the new cadence is not blocked by a stale prompt.

4.5 Consumption (hooks/useHeatGuard.ts:106-115)

```
if (autopilotSession is active within 4h)
  recheckIntervalMin = autopilot.intervalMin
  recheckUrgency     = autopilot.urgency
  autopilotActive    = true
else
  recheckIntervalMin = heatBandDefault
  recheckUrgency     = heatBandUrgency
  autopilotActive    = false
```

`components/RiskTimerDisplay.tsx` renders the resulting `timerSeconds` and shows an "URGENT" badge when state is DEPLETED or remaining time is < 5 min.

4.6 Hand-off and graceful exit

After 4 h, `useHeatGuard` falls back to band-default cadence — there is no need to explicitly "exit" Autopilot. The 4 h horizon is the **Sawka 2007** post-exercise rehydration window and was chosen so the cadence aligns with the period during which sweat-driven sodium and fluid replacement is physiologically meaningful.

4.7 Sodium positioning carried through Autopilot

The Sweat Calculator that triggers Autopilot also feeds the Recovery Intelligence card with two derived values that anchor AForce's "25 mg sodium is intentional" positioning:

- `aforceSodiumTotalMg = aforceSticks × AFORCE_SODIUM_PER_UNIT_MG (=25)` — `sweatRateEngine.ts:128, 509`
- `sodiumGapMg = max(0, sodiumLossMg – aforceSodiumTotalMg)` — the *intentional* gap, framed as "covered by structured-water absorption + marine minerals"

These values are visible in the user-facing audit copy ("125 mg delivered by 5 units; gap 1 723 mg covered by structured-water absorption") and are the basis of the comparison-table claims (AForce 25 mg / Gatorade ~270 mg / LMNT 1 000 mg / Liquid IV ~500 mg).

5. Autoscan (HydroScan)

Maps to: FIG. 10 (Autopilot / Autoscan) **Primary screen:** `artifacts/aforce-`

`os/screens/HydrationScanScreen.tsx` **Camera:** `components/CameraScanModal.tsx` **Service layer:** `services/hydrationScanService.ts`, `services/productRecognitionService.ts`, `services/comparisonEngine.ts`

Catalog: `data/productDatabase.ts`

5.1 What Autoscan is

A "scan-to-decide" pipeline. The user points the phone at any beverage barcode (or types one in). The system identifies the product, scores it against the user's *current physiological state* using the same engine described in §2/§3, and — if an AForce SKU would outperform the scanned product — issues a one-tap

replacement recommendation.

5.2 Input modes

Mode	Used on	Mechanism
Native barcode / QR	iOS, Android	expo-camera CameraView (CameraScanModal.tsx:38-48) reading ean13 , upc_a , qr , code128 . AForce-specific QR codes use the scheme aforce://product/<id>
Manual search	All platforms	Free-text product / brand / digit lookup
Mock scan tray	Web preview, demo, testing	HydrationScanScreen.tsx:217-250 exposes a fixture tray of barcodes for one-tap simulation

5.3 Resolution pipeline (services/productRecognitionService.ts)

1. Look the barcode up in the local `BARCODE_INDEX` (`data/productDatabase.ts`) — instant, offline, zero network.
2. If miss, fall back to the **Open Food Facts API** (`productRecognitionService.ts:126`) so any US beverage can be resolved even when not in the local catalog.
3. Normalize the resolved product into the local `ScannedProduct` schema (name, brand, category, hydration markers, electrolyte density, sugar level, mineral content).

5.4 Scoring (services/hydrationScanService.ts:55-63)

The product is scored on:

- **Hydration efficiency** (0-100 %) — mineral content + water content – sugar penalty
- **Fit score** (0-100) and **verdict** (Optimal / Strong / Acceptable / Suboptimal / Avoid) based on the user's current protocol context (Heat Stress, Recovery, Social, etc.) — i.e. the same `UserState` that drives §2/§3.

5.5 AForce comparison (hydrationScanService.ts:70 , 93-101)

The catalog is filtered to `isAForce: true` SKUs. The engine picks the best AForce product for the user's current state, computes its fit score, and compares to the scanned product. If `aforceFitScore ≥ scannedFitScore + 4` , an **AForce Replacement Card** (`components/AForceReplacementCard.tsx`) is rendered on the result, including the situational headline, detail, and a specific command ("Take 1 AForce Stick now with 16 oz water").

5.6 Output (ScanOutcome)

- `product` — resolved details
- `efficiency` — 0-100 % hydration efficiency
- `fitScore` + `verdict` — for the scanned product
- `aforceUplift` — if applicable: { `aforceProduct`, `fitDelta`, `command` }
- `recommendation` — situational headline + detail + command, rendered by `components/ScanResultCard.tsx`

5.7 Why Autoscan is novel

Existing barcode-scanning beverage apps return *static* nutrition facts. Autoscan scores a product against a *live* user state derived from the same Deterministic AI Layer that drives the rest of the app, and replaces a static "facts" output with a contextual decision (drink it / don't / use AForce instead). The replacement decision is auditable because both fit scores are computed by the same pure scoring code.

6. Patent figure ↔ implementation cross-reference

Figure	Title	Maps to section	Primary implementation file
FIG. 1	System Architecture	§1	artifacts/aforce-os/ , artifacts/api-server/
FIG. 2	Hero Intake Flow	§1 (home)	app/(tabs)/index.tsx , components/StatusPulseOrb.tsx
FIG. 3	Paywall Flow	§1 (entitlement)	hooks/useEntitlement.ts , app/subscription.tsx
FIG. 4	Social Mode Flow	§3.3	scoringEngine.ts:458-531 , services/bacEstimationService.ts
FIG. 5	Voice Command Flow	§1.4 (voice)	services/voiceService.ts , voiceTemplateEngine.ts
FIG. 6	Wireframes	§1.5	live iframes on canvas + screens/*
FIG. 7	Mobile App Architecture	§1	app/_layout.tsx , store/useAppStore.tsx
FIG. 8	Scoring Engine	§2	utils/scoringEngine.ts , services/hydrationScoreService.ts
FIG. 9	Deterministic AI Layer	§3	utils/scoringEngine.ts:537-591 , :728-748 , :790-838
FIG. 10	Autopilot / Autoscan	§4, §5	services/sweatRateEngine.ts:489-493 , services/hydrationScanService.ts

7. Suggested claim hooks (engineering input only — not legal advice)

Counsel may want to consider these implementation details as candidate claim elements:

1. **Pure-function pipeline computing a per-minute decaying performance score** from a defined set of biometric, environmental, and behavioral inputs, with a server- and client-side identity invariant (same input $\Rightarrow$ same output).
2. **Sweat-deficit-driven cadence override** that replaces a band-based recheck timer with a physiologically-derived interval ladder (8/12/20 min) for a bounded recovery window (4 h) following a measured sweat session.
3. **Beverage-recognition pipeline that scores any scanned product against a live user state** (not a static nutrition database) and conditionally surfaces a brand-specific replacement when the brand's fit score exceeds the scanned product's fit score by a configurable margin.
4. **Audit-on-output:** every score returned to the UI carries a `breakdown[]` array of bounded contributions so the user can inspect why the score is what it is.
5. **Mode-priority pipeline** in which a Social Mode ruleset (BAC + impairment) preempts the standard performance ladder and emits safety-class commands that the rest of the engine cannot override.

8. Verification

- TypeScript: `pnpm --filter @workspace/aforce-os typecheck` is clean as of this brief.
- Visual: AForce OS app preview, Recovery Card (Unlocked/Locked), and Hydration Score Band variants are live iframes on the engineering canvas; counsel can interact with them directly.
- Constants in this brief (25 mg sodium per AForce unit; 8 / 12 / 20 min autopilot ladder; 90 / 75 / 60 band thresholds) were re-verified by ripgrep against the source files cited above on the date of this brief.

End of main brief. Appendix A follows.

Appendix A — Patent Figures (FIG. 1 - FIG. 10)

This appendix reproduces all ten patent figures in filing order. Each figure is also embedded inline in the brief at the section that references it. Figures are also available as standalone vector files (fig{1..10}-*.svg) and as a separate "Figures" PDF (aforce-os-figures.pdf) for the firm's filing package.

A.1 — FIG. 1: System Architecture (maps to §1)





FIG. 1 — System Architecture. Top-level diagram of the AForce OS mobile client, the api-server, the Postgres / Drizzle data store, and the WebSocket realtime channel.

A.2 — FIG. 2: Hero Intake Flow (maps to §1.1, home screen)

FIG. 2 - Hero Intake and Score Recompute Cycle

User logs an AForce intake; orb pulse, recheck timer, and recovery banner update in lockstep




AForce OS - Patent Figure

FIG. 2 — Hero Intake Flow. The Hydration Control Center loop: status orb → AI command → user logs intake → score recompute → orb re-renders.

A.3 — FIG. 3: Paywall Flow (maps to §1.4, entitlement)

FIG. 3 — Paywall: Recovery Mode Unlock via Stripe

Locked card → checkout session → webhook-driven entitlement → live unlock



AForce OS - Patent Figure

FIG. 3 — Paywall Flow. Premium gating for Recovery Card, Sweat Calculator, Voice, Heat Guard, and Autoscan via `useEntitlement`.

A.4 — FIG. 4: Social Mode Flow (maps to §3.3)




FIG. 4 — Social Mode Flow. Drink-logging → Widmark BAC estimate → impairment level → safety-class command override.

A.5 — FIG. 5: Voice Command Flow (maps to §1.4, voice)




FIG. 5 — Voice Command Flow. Audio transcript → intent classification (LOG_INTAKE, GET_STATUS, ...) → action dispatch → spoken response via the brand template engine.

A.6 — FIG. 6: Wireframes (maps to \$1.5)




FIG. 6 — Wireframes. Hydration Control Center, Performance Signals (check), Recovery Protocol, Store, and Profile screens.

A.7 — FIG. 7: Mobile App Architecture (maps to §1.2)




FIG. 7 — Mobile App Architecture. Sliced React Context + pure reducer + memoized projections; expo-router file-based navigation; Clerk auth bridge.

A.8 — FIG. 8: Scoring Engine (maps to §2)




FIG. 8 — Scoring Engine. Pure-function pipeline from UserState snapshot → bounded contribution sum → score / band / breakdown / prediction / command.

A.9 — FIG. 9: Deterministic AI Layer (maps to §3)




FIG. 9 — Deterministic AI Layer. Mode classification → penalty application → bonus application → decay projection → command generation → recheck cadence.

A.10 — FIG. 10: Autopilot / Autoscan (maps to §4 and §5)




FIG. 10 — Autopilot / Autoscan. Sweat session → deficit % → autopilot cadence ladder (8 / 12 / 20 min, 4 h window). Barcode scan → product recognition → live-state fit score → AForce replacement decision.

End of brief. For follow-up — additional cross-references, formula derivations, or extracts of specific source files — contact AForce engineering.